BRYAN TALAVERA

Data Analyst · Analytics Engineer | UCSD B.S. Data Science

San Diego, CA · bryantalavera.gomez@gmail.com · github.com/Bryan-Tal · bryan-tal.github.io · linkedin.com/in/bryan-t163001290

SUMMARY

UCSD B.S. Data Science graduate building production-grade analytics and data engineering work, including a multi-phase NYC TLC ELT pipeline (Airflow, dbt, Snowflake) and tuned ML classifiers (0.95 AUC churn, 83% F1 taxi gratuities). Proficient in SQL, Python, Tableau, and the modern data stack. Targeting full-time data analyst, analytics engineer, or data engineer roles.

EDUCATION

University of California, San Diego · B.S. Data Science, Minor in Cognitive Science 2024

CERTIFICATIONS

Google Advanced Data Analytics Professional Certificate · Coursera · 2024 · View credential

PROJECTS

NYC Taxi ELT Pipeline | *Python · Airflow · dbt · Snowflake · AWS S3 · Docker · GitHub Actions* *2026*
Phase 1 of 6 complete · github.com/Bryan-Tal/nyc-taxi-elt

- Designed a batch ELT pipeline for 100M+ NYC TLC trip records into Snowflake, using a Medallion architecture (Raw → Staging → Marts) modeled in dbt with star-schema dimensional design
- Secured cross-cloud authentication via AWS IAM role assumption and least-privilege Snowflake RBAC, separating functional and access roles
- Containerized the orchestration layer with Docker and configured GitHub Actions CI/CD for repeatable, idempotent deployments

Predicting Employee Churn | *Python · Scikit-learn · XGBoost · Pandas · GridSearchCV* *2025*
github.com/Bryan-Tal/Predicting_Employee_Churn

- Compared Decision Tree, Random Forest, and XGBoost classifiers on an 11,991-row HR dataset; selected Random Forest as champion via 5-fold GridSearchCV refit on ROC-AUC
- Achieved 0.95 AUC, 97.97% accuracy, and 90% recall on the held-out test set; engineered an overworked feature (>175 hrs/mo) that ranked among the top predictors
- Surfaced that 100% of employees with 7 concurrent projects left the company, identifying a workload threshold for HR retention strategies

Predicting NYC Taxi Gratuities | *Python · XGBoost · Scikit-learn · Pandas · GridSearchCV* *2025*
github.com/Bryan-Tal/Predicting_NYC_Taxi_Gratuities

- Built a binary classifier predicting generous tippers (>20%) on 408,000+ NYC yellow taxi trips, tuning XGBoost via 5-fold GridSearchCV refit on F1
- Achieved 83.3% accuracy and 84% F1 on the held-out test set, beating the Random Forest baseline by +12.7 F1 points
- Engineered time-of-day, day-of-week, and rate-code features from raw timestamps; documented cash-payment survivorship bias as a model limitation

Predicting House Prices with OLS & Ridge Regression | *Python · NumPy · SciPy · Pandas · Matplotlib* *2025*
github.com/Bryan-Tal/Predicting_House_Prices_OLS

- Implemented Multiple Linear Regression and Ridge Regression from scratch using the closed-form normal equation in NumPy, foregoing scikit-learn for the regression itself
- Validated all 5 OLS assumptions (Shapiro-Wilk, Anderson-Darling, K-S, VIF) and searched 4 scaling methods × 7 Ridge α values; champion model achieved $R^2 = 0.66$, RMSE = 0.108

Med-Dash · UC San Diego Capstone | *React · MongoDB · ECharts · Vital API* *Sep 2023 to Mar 2024*

- Built a personalized, interactive dashboard for patients to visualize 3 daily health metrics (heart rate, sleep, activity) ingested from the Vital API
- Collaborated with a 4-person interdisciplinary team to integrate data persistence (MongoDB) and frontend visualization (React + ECharts); led validation across 3 wearable sensor types

SKILLS

Languages: Python, SQL, JavaScript, R
Machine Learning & Statistics: Scikit-learn, XGBoost, Pandas, NumPy, SciPy, Statistical Testing (Shapiro-Wilk, Anderson-Darling, K-S, VIF), GridSearchCV
Data Engineering: Apache Airflow, dbt, Snowflake, AWS S3 & IAM, Docker, GitHub Actions
Visualization & BI: Tableau, Matplotlib, Seaborn, ECharts, D3.js
Tools & Databases: Git/GitHub, Jupyter, MongoDB, PostgreSQL, MS SQL Server

WORK EXPERIENCE

In-Store Shopper, Pavilions · San Diego, CA April 2025 - Present

- Fulfilled 20+ time-sensitive grocery orders per shift while maintaining a 98%+ on-time completion rate
- Trained new associates on workflow procedures and order-fulfillment best practices, reducing handoff delays during shift transitions